

Mod_cluster doc

1. Overview	1
1.1. Platforms	1
1.2. Advantages	1
1.3. Requirements	2
1.4. Downloads	2
1.5. Quick Start	3
1.6. Configuration	3
1.7. Migration	3
1.8. SSL support	3
1.9. Load Balancing Demo Application	3
1.10. Strong cryptography warning	3
2. Quick Start Guide	5
3. native configuration	13
3.1. Apache httpd configuration	13
3.2. mod_proxy configuration	13
3.3. mod_slotmem configuration	13
3.4. mod_proxy_cluster	14
3.5. mod_manager	15
3.6. mod_advertise	18
3.7. Minimal Example	18
4. Building mod_cluster from sources	21
4.1. Build a patched httpd from it sources	21
4.2. Build the 4 modules of mod_cluster	21
4.3. Build the mod_proxy module	22
5. Installation of the httpd part	23
5.1. Configuration	23
5.2. Installing and using the bundles	23
6. java configuration	25
6.1. JBoss AS	25
6.1.1. Lifecycle listener	25
6.1.2. Connectors	26
6.1.3. Node Identity	26
6.1.4. Non-Clustered Mode	26
6.1.5. Clustered Mode	27
6.2. JBoss Web & Tomcat	28
6.2.1. Lifecycle Listener	28
6.2.2. Additional Tomcat dependencies	29
7. Building mod_cluster's Java components from source	31
7.1. Build products	31
8. java properties	33
8.1. Configuration Properties	33
8.1.1. Proxy Discovery Configuration	33
8.1.2. Proxy Configuration	35
8.1.3. SSL Configuration	37

8.1.4. HA Configuration	38
8.1.5. Load Configuration	38
9. Load Metrics	39
9.1. Server-Side Load Metrics	39
9.2. Web Container metrics	40
9.3. System/JVM metrics	44
9.4. Other metrics	45
10. Installation of the java part	47
10.1. Install the Java-side binaries	47
10.1.1. Installing in JBoss AS 6.0.0.M1 and up	47
10.1.2. Installing in JBoss AS 5.x	47
10.1.3. Installing in JBoss Web or Tomcat	47
10.1.4. Installing in JBoss AS 4.2.x or 4.3.x	47
11. Using https between httpd and JBossWEB	49
12. Using SSL in mod_cluster	51
12.1. Using SSL between JBossWEB and httpd	51
12.1.1. Apache httpd configuration part	51
12.1.2. ClusterListener configuration part	51
12.1.3. mod-cluster-jboss-beans configuration part	52
12.1.4. How the diferent files were created	52
12.2. Using SSL between httpd and JBossWEB	54
12.3. Forwarding SSL browser informations when using http/https between httpd and JBossWEB	54
13. Migration from mod_jk	57
14. Migration from mod_proxy	59
14.1. Workers	59
14.2. Balancers	60
15. Load Balancing Demo Application	63
15.1. Overview	63
15.2. Basic Usage	63
15.3. Client Driver Configuration Options	66
15.4. Load Generation Scenarios	67
16. Change Log	71
16.1. 1.0.10 (14 April 2011)	71
16.2. 1.0.9 (17 March 2011)	71
16.3. 1.0.8 (15 Feb 2011)	72
16.4. 1.0.7 (not released)	72
16.5. 1.0.6 (05 January 2011)	72
16.6. 1.0.5 (23 November 2010)	72
16.7. 1.0.4 (27 July 2010)	73
16.8. 1.0.3.GA (15 January 2010)	73
16.9. 1.0.2.GA (20 August 2009)	74
16.10. 1.0.1.GA (29 July 2009)	75
16.11. 1.0.0.GA (29 May 2009)	75

16.12. 1.0.0.CR2 (29 April 2009)	76
16.13. 1.0.0.CR1 (24 March 2009)	76
16.14. 1.0.0.Beta4 (14 Feb 2009)	77
16.15. 1.0.0.Beta3 (31 Jan 2009)	78
16.16. 1.0.0.Beta2 (12 Dec 2008)	78
16.17. 1.0.0.Beta1 (05 Nov 2008)	79
17. Frequently Asked questions	81
17.1. What is Advertise	81
17.2. What to do if I don't want to use Advertise (multicast):	81
17.3. I am using Tomcat 7 / 6 what should I do:	82

Overview

mod_cluster is an httpd-based load balancer. Like mod_jk and mod_proxy, mod_cluster uses a communication channel to forward requests from httpd to one of a set of application server nodes. Unlike mod_jk and mod_proxy, mod_cluster leverages an additional connection between the application server nodes and httpd. The application server nodes use this connection to transmit server-side load balance factors and lifecycle events back to httpd via a custom set of HTTP methods, affectionately called the Mod-Cluster Management Protocol (MCMP). This additional feedback channel allows mod_cluster to offer a level of intelligence and granularity not found in other load balancing solutions.

Within httpd, mod_cluster is implemented as a set of modules for httpd with mod_proxy enabled. Much of the logic comes from mod_proxy, e.g. mod_proxy_ajp provides all the AJP logic needed by mod_cluster.

1.1. Platforms

JBoss already prepares *binary packages* [http://www.jboss.org/mod_cluster/downloads.html] with httpd and mod_cluster so you can quickly try mod_cluster on the following platforms:

1. Linux x86, x64, ia64
2. Solaris x86, SPARC
3. Windows x86, x64, ia64
4. HP-UX PA-RISC, ia64

1.2. Advantages

mod_cluster boasts the following advantages over other httpd-based load balancers:

Dynamic configuration of httpd workers

Traditional httpd-based load balancers require explicit configuration of the workers available to a proxy. In mod_cluster, the bulk of the proxy's configuration resides on the application servers. The set of proxies to which an application server will communicate is determined either by a static list or using dynamic discovery via the *advertise* mechanism. The application server relays lifecycle events (e.g. server startup/shutdown) to the proxies allowing them to effectively auto-configure themselves. Notably, the graceful shutdown of a server will not result in a failover response by a proxy, as is the case with traditional httpd-based load balancers.

Server-side load balance factor calculation

In contrast with traditional httpd-based load balancers, mod_cluster uses load balance factors calculated and provided by the application servers, rather than computing these in the proxy.

Consequently, `mod_cluster` offers a more robust and accurate set of load metrics than is available from the proxy. (see [Load Metrics](#) for more)

Fine grained web-app lifecycle control

Traditional `httpd`-based load balancers do not handle web application undeployments particularly well. From the proxy's perspective requests to an undeployed web application are indistinguishable from a request for a non-existent resource, and will result in 404 errors. In `mod_cluster`, each server forwards any web application context lifecycle events (e.g. web-app deploy/undeploy) to the proxy informing it to start/stop routing requests for a given context to that server.

AJP is optional

Unlike `mod_jk`, `mod_cluster` does not require AJP. `httpd` connections to application server nodes can use HTTP, HTTPS, or AJP.

The original concepts are described in a [wiki](http://www.jboss.org/community/docs/DOC-11431) [<http://www.jboss.org/community/docs/DOC-11431>].

1.3. Requirements

1. `httpd-2.2.8+`
2. JBoss AS 5.0.0+ or JBossWeb 2.1.1+

Note: `httpd-2.2.8+` is already in the bundles, so if you use the bundle you don't need to download Apache `httpd`.

1.4. Downloads

Download the latest `mod_cluster` release [here](#) [http://www.jboss.org/mod_cluster/downloads/latest].

The release is comprised of the following artifacts:

1. `httpd` binaries for common platforms
2. JBoss AS/JBossWeb/Tomcat binary distribution

Alternatively, you can build from source using the subversion repository:

http://anonsvn.jboss.org/repos/mod_cluster/tags/ [release](#) [http://anonsvn.jboss.org/repos/mod_cluster/tags/] / [http://anonsvn.jboss.org/repos/mod_cluster/tags/]

1. Building [httpd modules](#)
2. Building [Java archives](#)

1.5. Quick Start

If you want to skip the details and just set up a minimal working installation of mod_cluster, see the [Quick Start Guide](#).

1.6. Configuration

1. Configuring [httpd](#)
2. Configuring [JBoss AS/Web](#)
3. Configuring [Tomcat](#)

1.7. Migration

Migrating from mod_jk or mod_proxy is fairly straight forward. In general, much of the configuration previously found in httpd.conf is now defined in the application server nodes.

1. Migrating from [mod_jk](#)
2. Migrating from [mod_proxy](#)

1.8. SSL support

Both the request connections between httpd and the application server nodes, and the feedback channel between the nodes and httpd can be secured. The former is achieved via the [mod_proxy_https module](#) and a corresponding ssl-enabled HTTP connector in JBoss Web. The latter requires the [mod_ssl module](#) and [explicit configuration in JBoss AS/Web](#).

1.9. Load Balancing Demo Application

The mod_cluster binary distribution for JBoss AS/JBossWeb/Tomcat includes a [demo application](#) that helps demonstrate how different server-side scenarios affect the routing of client requests by the load balancer. The demo application is located in the mod_cluster distribution's demo directory.

1.10. Strong cryptography warning

Mod_cluster contains mod_ssl, therefore the warning (copied from [OpenSSL](#) [<http://www.openssl.org/>])

PLEASE REMEMBER THAT EXPORT/IMPORT AND/OR USE OF STRONG CRYPTOGRAPHY SOFTWARE, PROVIDING CRYPTOGRAPHY HOOKS OR EVEN JUST COMMUNICATING TECHNICAL DETAILS ABOUT CRYPTOGRAPHY SOFTWARE IS ILLEGAL IN SOME PARTS OF THE WORLD. SO, WHEN YOU IMPORT THIS PACKAGE TO YOUR COUNTRY, RE-

DISTRIBUTE IT FROM THERE OR EVEN JUST EMAIL TECHNICAL SUGGESTIONS OR EVEN SOURCE PATCHES TO THE AUTHOR OR OTHER PEOPLE YOU ARE STRONGLY ADVISED TO PAY CLOSE ATTENTION TO ANY EXPORT/IMPORT AND/OR USE LAWS WHICH APPLY TO YOU. THE AUTHORS OF OPENSSEAL ARE NOT LIABLE FOR ANY VIOLATIONS YOU MAKE HERE. SO BE CAREFUL, IT IS YOUR RESPONSIBILITY.

Quick Start Guide

The following are the steps to set up a minimal working installation of mod_cluster on a single httpd server and a single back end server, either JBoss AS, JBossWeb or Tomcat. The steps can be repeated to add as many httpd servers or back end servers to your cluster as is desired.

The steps shown here are not intended to demonstrate how to set up a production install of mod_cluster; for example [using SSL to secure access](#) to the httpd-side mod_manager component is not covered. See the [httpd-side](#) and [java-side](#) configuration documentation for the full set of configuration options.

2.1. Download mod_cluster components

Download the latest [httpd and java release bundles](http://www.jboss.org/mod_cluster/downloads/latest.html) [http://www.jboss.org/mod_cluster/downloads/latest.html]. If there is no pre-built httpd bundle appropriate for your OS or system architecture, you can [build the binary from source](#).

2.2. Install the httpd binary

2.2.1. Install the whole httpd

The httpd-side bundles are gzipped tars and include a full httpd install. As they contain already an Apache httpd install you don't need to download Apache httpd. Just extract them in root, e.g.

```
cd /  
tar xvf mod-cluster-1.0.0-linux2-x86-ssl.tar.gz
```

That will give you a full httpd install in your /opt/jboss directory.

2.2.2. Install only the modules

If you already have a working httpd install that you would prefer to use, you'll need to download the bundle named mod_cluster httpd dynamic libraries corresponding to you plaform, extract the modules and copy them directory to your httpd install's module directory.

```
cd /tmp  
tar xvf mod_cluster-1.0.0.CR2-linux2-x86-so.tar.gz
```

And then you have to copy the files below to you module directory:

- * mod_proxy.so

- * mod_proxy_ajp.so
- * mod_slotmem.so
- * mod_manager.so
- * mod_proxy_cluster.so
- * mod_advertise.so

2.2.3. Install in home directory

Since 1.1.0.CR2 a script `opt/jboss/httpd/sbin/installhome.sh` allows reconfiguration of the bundle installation so that it can run in user's home directory. To do that just extract the bundle in your home directory and run the script. Once that done, httpd will run on port 8000 and will accept MCMP messages on `localhost:6666` and offer `/mod_cluster_manager` on the same host and port.

2.2.4. Install in windows

To install in windows unzip the bundle corresponding to your architecture. You may run httpd directly by using:

```
httpd-2.2\bin\httpd.exe -k start
```

or install Apache httpd as a service:

```
httpd-2.2\bin\httpd.exe -k install
```

and start the service:

```
net start Apache2.2
```

Note that in the windows bundles have a flat directory structure, so you have `httpd-2.2/modules/` instead of `opt/jboss/httpd/lib/httpd/modules`.

2.3. Configure httpd

Since 1.1.0.CR2 `httpd.conf` is preconfigured with the Quick Start values. You should adapt the default values to your configuration with older `mod_cluster` we will have to add the following to `httpd.conf`. If you extracted the download bundle to root as shown above and are using that extract as your httpd install, `httpd.conf` is located in `/opt/jboss/httpd/httpd/conf`.

```
LoadModule proxy_module /opt/jboss/httpd/lib/httpd/modules/mod_proxy.so
LoadModule proxy_ajp_module /opt/jboss/httpd/lib/httpd/modules/mod_proxy_ajp.so
LoadModule slotmem_module /opt/jboss/httpd/lib/httpd/modules/mod_slotmem.so
LoadModule manager_module /opt/jboss/httpd/lib/httpd/modules/mod_manager.so
LoadModule proxy_cluster_module /opt/jboss/httpd/lib/httpd/modules/mod_proxy_cluster.so
LoadModule advertise_module /opt/jboss/httpd/lib/httpd/modules/mod_advertise.so

Listen 10.33.144.3:6666
<VirtualHost 10.33.144.3:6666>

    <Directory />
        Order deny,allow
        Deny from all
        Allow from 10.33.144.
    </Directory>

    KeepAliveTimeout 60
    MaxKeepAliveRequests 0

    ManagerBalancerName mycluster
    AdvertiseFrequency 5

</VirtualHost>
```

If you are using your own install of httpd, httpd.conf is found in your install's conf directory. The content to add to httpd.conf is slightly different from the above (different path to the various .so files):

```
LoadModule proxy_module modules/mod_proxy.so
LoadModule proxy_ajp_module modules/mod_proxy_ajp.so
LoadModule slotmem_module modules/mod_slotmem.so
LoadModule manager_module modules/mod_manager.so
LoadModule proxy_cluster_module modules/mod_proxy_cluster.so
LoadModule advertise_module modules/mod_advertise.so

Listen 10.33.144.3:6666
<VirtualHost 10.33.144.3:6666>

    <Directory />
        Order deny,allow
        Deny from all
        Allow from 10.33.144.
```

```
</Directory>

KeepAliveTimeout 60
MaxKeepAliveRequests 0

ManagerBalancerName mycluster
AdvertiseFrequency 5

</VirtualHost>
```

2.4. Install the Java-side binaries

First, extract the java-side binary to a temporary directory. The following assumes it was extracted to /tmp/mod-cluster

Your next step depends on whether your target server is JBoss AS or JBossWeb/Tomcat.

2.4.1. Installing in JBoss AS 6.0.0.M1 and up

You don't need to do anything to install the java-side binaries in AS 6.x; it's part of the AS distribution's default, standard and all configurations.

2.4.2. Installing in JBoss AS 5.x

Assuming \$JBOSS_HOME indicates the root of your JBoss AS install and that you want to use mod_cluster in the AS's all config:

```
cp -r /tmp/mod-cluster/mod-cluster.sar $JBOSS_HOME/server/all/deploy
```

2.4.3. Installing in JBoss Web or Tomcat

Assuming \$CATALINA_HOME indicates the root of your JBossWeb or Tomcat install:

```
cp -r /tmp/mod-cluster/JBossWeb-Tomcat/lib/*
    $CATALINA_HOME/lib/
```

2.4.4. Installing in JBoss AS 4.2.x or 4.3.x

Assuming \$JBOSS_HOME indicates the root of your JBoss AS install and that you want to use mod_cluster in the AS's all config:

```
cp -r /tmp/mod-cluster/JBossWeb-Tomcat/lib/mod-cluster.jar $JBOSS_HOME/server/all/
deploy/jboss-web.deployer
```

2.5. Configuring JBoss AS, JBoss Web or Tomcat

mod_cluster is integrated into your web server by adding an implementation of the Tomcat LifecycleListener interface to the server's configuration, and by giving each node a unique name.

2.5.1. Integrate mod_cluster with JBoss AS's embedded JBossWeb

<listitem>

Edit the \$JBOSS_HOME/server/all/deploy/jbossweb.sar/server.xml file, adding the following next to the other Listener elements:

</listitem>

```
<Listener

className="org.jboss.web.tomcat.service.deployers.MicrocontainerIntegrationLifecycleListener"
delegateBeanName="ModClusterService"/>
```

2.5.2. Integrate mod_cluster with standalone JBossWeb or Tomcat

Edit the \$CATALINA_HOME/conf/server.xml file, adding the following next to the other Listener elements:

```
<Listener
className="org.jboss.modcluster.ModClusterListener"
advertise="true"/>
```

2.5.3. Integrate mod_cluster with JBoss AS 4.2.x and 4.3.x

Edit the \$JBOSS_HOME/server/all/deploy/jboss-web.deployer/server.xml file, adding the following next to the other Listener elements:

```
<Listener
className="org.jboss.modcluster.ModClusterListener"
```

```
advertise="true"/>
```

2.5.4. Configure a jvmRoute

Each back end server in your cluster needs to have a unique name, known as a "jvmRoute". This is configured by adding an attribute to the Engine element in server.xml.

```
<Engine name="jboss.web" defaultHost="localhost"  
  jvmRoute="node01">
```

If your cluster includes multiple back-end servers, you must assign a different jvmRoute to each server.

Beginning with the 1.1.0.BETA1 release, configuring a jvmRoute is no longer absolutely required; if one isn't provided mod_cluster will generate one from the address and port of the Connector used for receiving requests, plus the name of the Tomcat Engine. Still, configuring a jvmRoute is recommended, since the jvmRoute is appended to all session ids. The generated jvmRoute is lengthy and includes information you may not want to expose to the internet via session ids.

2.6. Start httpd

To start httpd do the following:

```
/opt/jboss/httpd/sbin/apachectl start
```

2.7. Start the back end server

2.7.1. Starting JBoss AS

```
cd $JBOSS_HOME/bin ./run.sh -c all
```

2.7.2. Starting JBossWeb or Tomcat

```
cd $CATALINA_HOME ./startup.sh
```


2.8. Set up more back end servers

Repeat the back end server install and configuration steps; just be sure to use a different `jvmRoute` for each server.

2.9. Experiment with the Load Balancing Demo Application

The [load balancing demo application](#) is a good way to learn about `mod_cluster`'s capabilities.

native configuration

3.1. Apache httpd configuration

You need to load the modules that are needed for mod_cluster for example:

```
LoadModule proxy_module modules/mod_proxy.so
LoadModule proxy_ajp_module modules/mod_proxy_ajp.so
LoadModule slotmem_module modules/mod_slotmem.so
LoadModule manager_module modules/mod_manager.so
LoadModule proxy_cluster_module modules/mod_proxy_cluster.so
LoadModule advertise_module modules/mod_advertise.so
```

mod_proxy and mod_proxy_ajp are standard httpd modules. mod_slotmem is a shared slotmem memory provider. mod_manager is the module that reads information from JBoss AS/JBossWeb/Tomcat and updates the shared memory information. mod_proxy_cluster is the module that contains the balancer for mod_proxy. mod_advertise is an additional module that allows httpd to advertise via multicast packets the IP and port where the mod_cluster is listening. This multi-module architecture allows the modules to easily be changed depending on what the customer wants to do. For example when using http instead of AJP, only

```
LoadModule proxy_ajp_module modules/mod_proxy_ajp.so
```

needs to be changed to:

```
LoadModule proxy_http_module modules/mod_proxy_http.so
```

3.2. mod_proxy configuration

mod_proxy directives like ProxyIOBufferSize could be used to configure mod_cluster. There is no need to use ProxyPass directives because mod_cluster automatically configures which URLs have to be forwarded to JBossWEB.

3.3. mod_slotmem configuration

The actual version doesn't require any configuration directives.

3.4. mod_proxy_cluster

CreateBalancers: Define how the balancer are created in the httpd VirtualHosts, this is to allow directives like:

```
ProxyPass / balancer://mycluster1/
```

0: Create in all VirtualHosts defined in httpd.

1: Don't create balancers (requires at least one ProxyPass/ProxyPassMatch to define the balancer names).

2: Create only the main server.

Default: 2

Note: When using 1 don't forget to configure the balancer in the ProxyPass directive, because the default is empty stickysession and nofailover=Off and the values received via the MCMP CONFIG message are ignored.

UseAlias: Check that the Alias corresponds to the ServerName (See [Host Name Aliases](http://labs.jboss.com/file-access/default/members/jbossweb/freezone/docs/latest/config/host.html) [http://labs.jboss.com/file-access/default/members/jbossweb/freezone/docs/latest/config/host.html])

0: Don't (ignore Aliases)

1: Check it

Default: 0 Ignore the Alias information from the nodes.

LBstatusRecalTime:

Time interval in seconds for loadbalancing logic to recalculate the status of a node: (Default: 5 seconds)

The actual formula to recalculate the status of a node is:

```
status = lbstatus + (elected - oldelected) * 1000)/lbfactor;
```

lbfactor is received for the node via STATUS messages.lbstatus is recalculated every LBstatusRecalTime seconds using the formula:

```
lbstatus = (elected - oldelected) * 1000)/lbfactor;
```

elected is the number of time the worker was elected. `oldelected` is elected last time the `lbstatus` was recalculated. The node with the lowest status is selected. Nodes with `lbfactor` # 0 are skipped by the both calculation logic.

ProxyPassMatch/ProxyPass: `ProxyPassMatch` and `ProxyPass` are `mod_proxy` directives that when using `!` (instead the back-end url) prevent to reverse-proxy in the path. This could be used allow `httpd` to serve static information like images.

```
ProxyPassMatch ^(/.*\.(gif))$ !
```

The above for example will allow `httpd` to server directly the `.gif` files.

3.5. mod_manager

The Context of a `mod_manger` directive is `VirtualHost` except mentioned otherwise. `server config` means that it must be outside a `VirtualHost` configuration. If not an error message will be displayed and `httpd` won't start.

MemManagerFile: That is the base name for the names `mod_manager` will use to store configuration, generate keys for shared memory or lock files. That must be an absolute path name; the directories will created if needed. It is highly recommended that those files are placed on a local drive and not an NFS share. (Context: server config) Default: `$server_root/logs/`

Maxcontext: That is the number max of contexts supported by `mod_cluster`. (Context: server config) Default: 100

Maxnode: That is the number max nodes supported by `mod_cluster`. (Context: server config) Default: 20

Maxhost: That is the number max host (Aliases) supported by `mod_cluster`. That is also the max number of balancers. (Context: server config) Default: 10

Maxsessionid: That is the number of active sessionid we store to give number of active sessions in the `mod_cluster-manager` handler. A session is unactive when `mod_cluster` doesn't receive any information from the session in 5 minutes. (Context: server config) Default: 0 (the logic is not activated).

ManagerBalancerName: That is the name of balancer to use when the JBoss AS/JBossWeb/Tomcat doesn't provide a balancer name. That is not supported in the current version. (Default value is used). Default: `mycluster`

PersistSlots: Tell `mod_slotmem` to persist the nodes, Alias and Context in files. (Context: server config) Default: Off

CheckNonce: Switch check of nonce when using mod_cluster-manager handler on | off (Default: on Nonce checked) Since 1.1.0.CR1

SetHandler mod_cluster-manager: That is the handler to display the node mod_cluster sees from the cluster. It displays the information about the nodes like INFO and additionally counts the number of active sessions.

```
<Location /mod_cluster-manager>  
SetHandler mod_cluster-manager  
Order deny,allow  
Deny from all  
Allow from 127.0.0.1  
</Location>
```

When accessing the location you define in httpd.conf you get something like:

Mod_cluster Status - Mozilla Firefox

View History Bookmarks Tools Help

http://localhost:8000/mod_cluster-manager

Google

Smart Bookmarks Red Hat Magazine Red Hat -- Intranet ... The World Clock - T... Red Hat -- HolidayS... EMEA Service - Liv...

fresh

jvm2 (ajp://127.0.0.1:8099): [Enable Contexts](#) [Disable Contexts](#)

mycluster,Domain: ,Flushpackets: Off,Flushwait: 10000,Ping: 10000000,Swax: 20,Ttl: 60000000,Elected: 0,Read: 0,Transferred: 0,Connected: 0,Load: 100,Num sessions

al Host 1:

Status: ENABLED [Disable](#)

, Status: ENABLED [Disable](#)

nger, Status: ENABLED [Disable](#)

status: ENABLED [Disable](#)

jvm1 (ajp://127.0.0.1:8009): [Enable Contexts](#) [Disable Contexts](#)

mycluster,Domain: ,Flushpackets: Off,Flushwait: 10000,Ping: 10000000,Swax: 20,Ttl: 60000000,Elected: 8,Read: 6384,Transferred: 46,Connected: 6,Load: 100,Num sessions

al Host 1:

Status: ENABLED [Disable](#)

atus: ENABLED [Disable](#)

nger, Status: ENABLED [Disable](#)

atus: ENABLED [Disable](#)

SessionIDs:

0148A1D6580E8C4B00E5AB766A, jvm1 route: jvm1

Note that:

Transferred: Corresponds to the POST data send to the back-end server.

Connected: Corresponds to the number of requests been processed when the mod_cluster status page was requested.

sessions: Corresponds to the number of sessions mod_cluster report as active (on which there was a request during the past 5 minutes). That field is not present when Maxsessionid is zero.

3.6. mod_advertise

mod_advertise uses multicast packets to advertise the VirtualHost where it is configured that must be the same VirtualHost where mod_manager is defined. Of course at least one mod_advertise must be in the VirtualHost to allow mod_cluster to find the right IP and port to give to the ClusterListener.

ServerAdvertise On: Use the advertise mechanism to tell the JBoss AS/JBossWeb/Tomcat to whom it should send the cluster information.

ServerAdvertise On http://hostname:port: Tell the hostname and port to use. Only needed if the VirtualHost is not defined correctly, if the VirtualHost is a *Name-based Virtual Host* [<http://httpd.apache.org/docs/2.2/vhosts/name-based.html>] or when VirtualHost is not used.

ServerAdvertise Off: Don't use the advertise mechanism.

Default: Off. (Any Advertise directive in a VirtualHost sets it to On in the VirtualHost)

AdvertiseGroup IP:port: That is the multicast address to use (something like 232.0.0.2:8888 for example). IP should correspond to AdvertiseGroupAddress and port to AdvertisePort in the JBoss AS/JBossWeb/Tomcat configuration. Note that if JBoss AS is used and the -u startup switch is included in the AS startup command, the default AdvertiseGroupAddress is the value passed via the -u. Default: 224.0.1.105:23364. If port is missing the default port will be used: 23364.

AdvertiseFrequency seconds[.milliseconds]: Time between the multicast messages advertising the IP and port. Default: 10 Ten seconds.

AdvertiseSecurityKey value: key string to identify the mod_cluster in JBossWEB. Default: No default value. Information not sent.

AdvertiseManagerUrl value: Not used in this version (It is sent in the X-Manager-Url: value header). That is the URL that JBoss AS/JBossWeb/Tomcat should use to send information to mod_cluster. Default: No default value. Information not sent.

AdvertiseBindAddress IP:port: That is the address and port httpd is bind to send the multicast messages. This allow to specify an address on multi IP address boxes. Default: 0.0.0.0:23364

3.7. Minimal Example

```
LoadModule proxy_module modules/mod_proxy.so
LoadModule proxy_ajp_module modules/mod_proxy_ajp.so
LoadModule slotmem_module modules/mod_slotmem.so
LoadModule manager_module modules/mod_manager.so
LoadModule proxy_cluster_module modules/mod_proxy_cluster.so
LoadModule advertise_module modules/mod_advertise.so

Listen 10.33.144.3:6666
```



```
<VirtualHost 10.33.144.3:6666>
```

```
  <Location />
```

```
    Order deny,allow
```

```
    Deny from all
```

```
    Allow from 10.33.144.
```

```
  </Location>
```

```
KeepAliveTimeout 60
```

```
MaxKeepAliveRequests 0
```

```
ManagerBalancerName mycluster
```

```
ServerAdvertise on
```

```
</VirtualHost>
```

Building mod_cluster from sources

4.1. Build a patched httpd from it sources

To build httpd-2.2.x from its sources see [ASF httpd doc](http://httpd.apache.org/docs/2.2/install.html) [http://httpd.apache.org/docs/2.2/install.html]

If needed, patch the httpd-2.2.x sources with (The patch prevents long waiting time when the node IP can't be resolved that should not happen so you can skip the patch part if you don't want to rebuild httpd). [mod_proxy_ajp.patch](http://anonsvn.jboss.org/repos/mod_cluster/trunk/native/mod_proxy_cluster/mod_proxy_ajp.patch) [http://anonsvn.jboss.org/repos/mod_cluster/trunk/native/mod_proxy_cluster/mod_proxy_ajp.patch]

```
(cd modules/proxy
 patch -p0 < $location/mod_proxy_ajp.patch
)
```

Configure httpd with something like:

```
./configure --prefix=apache_installation_directory \
  --with-mpm=worker \
  --enable-mods-shared=most \
  --enable-maintainer-mode \
  --with-expat=builtin \
  --enable-ssl \
  --enable-proxy \
  --enable-proxy-http \
  --enable-proxy-ajp \
  --disable-proxy-balancer
```

Rebuild (make) and reinstall (make install) after that.

4.2. Build the 4 modules of mod_cluster

You need an httpd installation with mod_proxy (--enable-proxy) and ajp protocol (--enable-proxy-ajp) enabled and with dso enabled (--enable-so)

Download the mod_cluster sources:

```
svn co http://anonsvn.jboss.org/repos/mod_cluster/trunk/ mod_cluster
```

Build the mod_cluster modules components, for each subdirectory advertise, mod_manager, mod_proxy_cluster and mod_slotmem do something like:

```
sh buildconf
./configure --with-apxs=apxs_file_location
make
cp *.so apache_installation_directory/modules
```

Where apache_installation_directory is the location of an installed version of httpd-2-2.x.

NOTE: You can ignore the libtool message on most platform:

```
libtool: install: warning: remember to run `libtool --finish apache_installation_directory/modules'
```

Once that is done use [Apache httpd configuration](#) to configure mod_cluster.

4.3. Build the mod_proxy module

It is only needed for httpd-2.2.x where x < 11. Process like the other mod_cluster modules.

Installation of the httpd part

Several bundles are available at http://www.jboss.org/mod_cluster/downloads.html.

In case you can't find a prepared package of mod_cluster in the download area, it is possible to build mod_cluster for the sources. You need a distribution of httpd (at least 2.2.8) or (better) a source tarball of httpd and the sources of mod_cluster. [Building](#) explains how to build mod_cluster for its sources.

5.1. Configuration

A minimal configuration is needed in httpd (See [httpd.conf](#)). A listener must be added in in JBossWEB conf/server.xml (See [Configuring JBoss AS/Web](#)).

5.2. Installing and using the bundles

The bundles are tar.gz on POSIX platforms just extract them in root something like:

```
cd /  
tar xvf mod-cluster-1.0.0-linux2-x86-ssl.tar.gz
```

The httpd.conf is located in /opt/jboss/httpd/httpd/conf to quick test just add something like:

```
Listen 10.33.144.3:6666  
<VirtualHost 10.33.144.3:6666>  
  
    <Directory />  
        Order deny,allow  
        Deny from all  
        Allow from 10.33.144.  
    </Directory>  
  
    KeepAliveTimeout 60  
    MaxKeepAliveRequests 0  
  
    ManagerBalancerName mycluster  
    AdvertiseFrequency 5  
  
</VirtualHost>
```

To start httpd do the following:

```
/opt/jboss/httpd/sbin/apachectl start
```

NOTE: Make sure to use SSL before going in production.

java configuration

6.1. JBoss AS

The entry point for configuring `mod_cluster` with JBoss AS is the `server.xml` file within JBoss Web service. This file is located at `$JBOSS_HOME/server/profile/deploy/jbossweb.sar/server.xml`

6.1.1. Lifecycle listener

`mod_cluster` leverages a `<Listener>` to inform itself of web container lifecycle events, e.g. deploy/undeploy of a web application, startup/shutdown of the server, etc.

We use the generic `MicrocontainerIntegrationLifecycleListener` to delegate all lifecycle events to a bean deployed by the JBoss microcontainer.

`mod_cluster` 1.0.x is wired directly to the Catalina container API. Catalina lifecycle events are delegated directly to `mod_cluster`'s core service.

```
<Listener

    delegateBeanName="ModClusterService"/><!-- Non-clustered mode -->
<!--Listener

    delegateBeanName="HAModClusterService"/><!-- Clustered mode -->
```

To ensure that the delegate bean exists before the JBoss Web service starts, and that the bean does not get destroyed until the service shuts down, we must add an explicit microcontainer dependency to the JBoss Web service. To do this, update the `WebServer` bean contained in `$JBOSS_HOME/server/profile/deploy/jbossweb.sar/META-INF/jboss-beans.xml` with the appropriate dependency:

```
<bean name="WebServer" class="org.jboss.web.tomcat.service.deployers.TomcatService">
  <!-- ... -->
  <depends>ModClusterService</depends><!-- Non-clustered mode -->
  <!--depends>HAModClusterService</depends--><!-- Clustered mode -->
  <!-- ... -->
</bean>
```

For details of how to configure [ModClusterService](#) or [HAModClusterService](#) refer to the appropriate section.

6.1.2. Connectors

Like `mod_jk` and `mod_proxy_balancer`, `mod_cluster` requires a connector in your `server.xml` on which to forward web requests. `mod_cluster` is not confined to AJP, but can use HTTP as well.

If multiple possible connectors are available, `mod_cluster` uses the following algorithm to choose between them:

1. If an native (APR) AJP connector is available, use it
2. If an AJP connector is available, use it
3. Otherwise, choose the connector with the highest max threads.

e.g.

```
<Server>
<!-- ... -->
<Service name="jboss.web">
    <Connector protocol="HTTP/1.1" port="${jboss.web.http.port}"
address="${jboss.bind.address}" redirectPort="${jboss.web.https.port}"/>
    <Connector protocol="AJP/1.3" port="${jboss.web.ajp.port}" address="${jboss.bind.address}"
redirectPort="${jboss.web.https.port}"/>
    <Engine name="jboss.web" defaultHost="localhost">
        <!-- ... -->
    </Engine>
</Service>
</Server>
```

6.1.3. Node Identity

Like `mod_jk` and `mod_proxy_balancer`, `mod_cluster` identifies each node via a unique `jvm` route. This is specified using the [jvmRoute](#) [<http://tomcat.apache.org/tomcat-6.0-doc/config/engine.html#Common%20Attributes>] attribute of the Engine within `server.xml`.

Users must define `jvm` routes manually, perhaps using system property expansion, e.g. `jvmRoute="${jboss.mod_cluster.jvmRoute}"`

6.1.4. Non-Clustered Mode

`mod_cluster` can operate in 2 modes: non-clustered or [clustered](#). In non-clustered mode, each JBoss server node communicates with a `httpd` proxy directly, and do not communicate with

each other. Non-clustered mode is powered by the ModClusterService, whose definition varies depending on the version of mod_cluster in use.

In general, ModClusterService defines:

1. A set of *configuration properties*
2. A mechanism for *providing the load balance factor* to a proxy

In version 1.0.x, all mod_cluster configuration properties are members of the ModClusterService bean itself:

```
<bean name="ModClusterService" class="org.jboss.modcluster.ModClusterService" mode="On
Demand">
  <constructor>
    <parameter class="org.jboss.modcluster.load.LoadBalanceFactorProvider">
      <inject bean="DynamicLoadBalanceFactorProvider"/>
    </parameter>
  </constructor>
  <!-- Configuration properties go here -->
  <property name="advertise">true</property>
</bean>
```

6.1.5. Clustered Mode

mod_cluster can operate in 2 modes: *non-clustered* or clustered. In clustered mode, a single JBoss node (the HA singleton master) communicates with the httpd proxy on behalf of the other nodes in the cluster. This mode offers the following advantages over non-clustered mode:

1. The state of each proxy will be kept in sync on each node in the cluster.
2. Proxies will be proactively notified of view changes, most notably, crashed members, allowing a proxy to gracefully reconfigure itself to avoiding failover processing.
3. Adds domain management functionality, providing the ability to enable, disable, or gracefully stop all nodes sharing the same domain. A domain is a logical grouping of nodes, typically representing a replication "buddy group".

Clustered mode is instrumented using JBoss clustering technologies and is configured via the HAModClusterService bean, the definition of which varies slightly between mod_cluster versions.

In general, HAModClusterService defines:

1. A set of *configuration properties*
2. A mechanism for *providing the load balance factor* to a proxy

3. An HAPartition, the JBoss clustering group communication building block. The default HAPartition is defined in: `$JBoss_HOME/server/profile/deploy/cluster/hapartition-jboss-beans.xml`

```
<bean name="HAModClusterService" class="org.jboss.modcluster.ha.HAModClusterService"
mode="On Demand">

annotation>
  <constructor>
    <parameter class="org.jboss.ha.framework.interfaces.HAPartition">
      <inject bean="HAPartition"/>
    </parameter>
    <parameter class="org.jboss.modcluster.config.ha.HAModClusterConfig">
      <inject bean="HAModClusterConfig"/>
    </parameter>
    <parameter class="org.jboss.modcluster.load.LoadBalanceFactorProvider">
      <inject bean="DynamicLoadBalanceFactorProvider"/>
    </parameter>
    <parameter class="org.jboss.ha.framework.interfaces.HASingletonElectionPolicy">
      <bean class="org.jboss.ha.singleton.HASingletonElectionPolicySimple"/>
    </parameter>
  </constructor>
</bean>
<bean name="HAModClusterConfig"
class="org.jboss.modcluster.config.ha.HAModClusterConfig" mode="On Demand">
  <property name="advertise">true</property>
</bean>
```

6.2. JBoss Web & Tomcat

In the absence of the ability for complex configuration, `mod_cluster`'s entire configuration for JBoss Web or Tomcat resides entirely within `$CATALINA_HOME/conf/server.xml`.

This adds the following constraints to `mod_cluster`'s feature set:

1. Only non-clustered mode is supported
2. *Only one load metric* can be used in the load balance factor calculation

6.2.1. Lifecycle Listener

The entry point for JBoss Web and Tomcat configuration is the `ModClusterListener`. All `mod_cluster` *configuration properties* are defined as attributes of the Listener element.

```
<Listener className="org.jboss.modcluster.ModClusterListener"
  advertise="true"/>
```

6.2.2. Additional Tomcat dependencies

mod_cluster uses jboss-logging, which exists in JBoss Web, but not in Tomcat. Consequently, to use mod_cluster with Tomcat, it is necessary to add [jboss-logging-spi.jar](http://repository.jboss.org/jboss/common-logging-spi/) [http://repository.jboss.org/jboss/common-logging-spi/] to \$CATALINA_HOME/lib.

Building mod_cluster's Java components from source

Building from source the components used in JBoss AS/JBossWeb/Tomcat requires Maven 2.0.x.

Steps to build:

1. Download the mod_cluster sources

```
svn co http://anonsvn.jboss.org/repos/mod_cluster/trunk/ mod_cluster
```

2. Use maven "dist" profile to build:

```
cd mod_cluster; mvn -P dist package
```

Note: Some maven tests are using UDP port 23365 make sure you are local firewall allows the port.

7.1. Build products

The build produces the following output in the target directory:

1. mod-cluster.sar

Exploded format sar to copy to the deploy dir in your JBoss AS install

2. JBossWeb-Tomcat/lib directory

Jar files to copy to the lib directory in your JBossWeb or Tomcat install to support use of mod_cluster

3. demo directory

The [load balancing demo](#) application

4. mod-cluster-XXX.tar.gz

The full distribution tarball; includes the above elements

java properties

8.1. Configuration Properties

The tables below enumerate the configuration properties available to an application server node. The location for these properties depends on [how mod_cluster is configured](#).

8.1.1. Proxy Discovery Configuration

The list of proxies from which an application expects to receive AJP connections is either defined statically, via the addresses defined in the [proxyList](#) configuration property; or discovered dynamically via the advertise mechanism. Using a special mod_advertise module, proxies can advertise their existence by periodically broadcasting a multicast message containing its address/port. This functionality is enabled via the [advertise](#) configuration property. If configured to listen, a server can learn of the proxy's existence, then notify that proxy of its own existence, and update its configuration accordingly. This frees both the proxy *and* the server from having to define static, environment-specific configuration values.

Attribute	Default	Description
proxyList	None	<p>Defines a comma delimited list of httpd proxies with which this node will initially communicate. Value should be of the form:</p> <p>address1:port1,address2:port2</p> <p>Using the default configuration, this property can be manipulated via the jboss.modcluster.proxyList system property.</p>
excludedContexts	ROOT,admin-console,invoker,jbossws,jmx-console,juddi,web-console	<p>List of contexts to exclude from httpd registration, of the form:</p> <p>host1:context1,host2:context2,host3:context3</p> <p>If no host is indicated, it is assumed to be the default host of the server (e.g. localhost). "ROOT" indicates the root context. Using the default configuration, this property can be manipulated</p>

Attribute	Default	Description
		via the <code>jboss.modcluster.excludedContexts</code> system property.
<code>autoEnableContexts</code>	<code>true</code>	If <code>false</code> the contexts are registered disabled in <code>httpd</code> , they need to be enabled via the <code>enable()</code> mbean method or via <code>mod_cluster_manager</code> . Since version 1.1.0.CR1.
<code>stopContextTimeout</code>	10	The number of seconds to wait for clean shutdown of a context (completion of pending requests for a distributable context; or destruction/expiration of active sessions for a non-distributable context). Since version 1.1.0.CR1.
<code>proxyURL</code>	<i>None</i>	If defined, this value will be prepended to the URL of MCMP commands.
<code>socketTimeout</code>	20000	Number of milliseconds to wait for a response from an <code>httpd</code> proxy to MCMP commands before timing out, and flagging the proxy as in error.
<code>advertise</code>	<code>true</code> , if <code>proxyList</code> is undefined, <code>false</code> otherwise	If enabled, <code>httpd</code> proxies will be auto-discovered via multicast announcements. This can be used either in concert or in place of a static <code>proxyList</code> .
<code>advertiseGroupAddress</code>	224.0.1.105	UDP address on which to listen for <code>httpd</code> proxy multicast advertisements
<code>advertisePort</code>	23364	UDP port on which to listen for <code>httpd</code> proxy multicast advertisements
<code>advertiseSecurityKey</code>	<i>None</i>	If specified, <code>httpd</code> proxy advertisements checksums will be verified using this value as a salt

8.1.2. Proxy Configuration

The following configuration values are sent to proxies during server startup, when a proxy is detected via the advertise mechanism, or during the resetting of a proxy's configuration during error recovery.

Attribute	Default	Description
stickySession	true	Indicates whether subsequent requests for a given session should be routed to the same node, if possible.
stickySessionRemove	false	Indicates whether the httpd proxy should remove session stickiness in the event that the balancer is unable to route a request to the node to which it is stuck. This property is ignored if stickySession is false.
stickySessionForce	true	Indicates whether the httpd proxy should return an error in the event that the balancer is unable to route a request to the node to which it is stuck. This property is ignored if stickySession is false.
workerTimeout	-1	Number of seconds to wait for a worker to become available to handle a request. When all the workers of a balancer are usable, mod_cluster will retry after a while (workerTimeout/100) to find an usable worker. That is timeout in the balancer mod_proxy documentation. A value of -1 indicates that the httpd will not wait for a worker to be available and will return an error if none is available.
maxAttempts	1	Number of times an httpd proxy will attempt to send a given request to a worker before giving up. The minimum

Attribute	Default	Description
		value is 1, meaning try only once. (Note that mod_proxy default is also 1: no retry).
flushPackets	false	Enables/disables packet flushing
flushWait	-1	Time to wait before flushing packets. A value of -1 means wait forever.
ping	10 seconds	Time to wait for a pong answer to a ping
smax	Determined by httpd configuration	Soft maximum idle connection count (that is the smax in worker mod_proxy documentation). The maximum value depends on the httpd thread configuration (ThreadsPerChild or 1).
ttl	60 seconds	Time to live (in seconds) for idle connections above smax
nodeTimeout	-1 (none)	Timeout (in seconds) for proxy connections to a node. That is the time mod_cluster will wait for the back-end response before returning error. That corresponds to timeout in the worker mod_proxy documentation. Note that mod_cluster always uses a cping/cpong before forwarding a request and the connectiontimeout value used by mod_cluster is the ping value.
balancer	mycluster	The balancer name
domain	None	If specified, load will be balanced among jvmRoutes with the same domain. This is primarily used in conjunction with partitioned

Attribute	Default	Description
		session replication (e.g. buddy replication).

Note: When `nodeTimeout` is not defined the `ProxyTimeout` directive `Proxy` is used. If `ProxyTimeout` is not defined the server timeout (`Timeout`) is used (default 300 seconds). `nodeTimeout`, `ProxyTimeout` or `Timeout` is set at the socket level.

8.1.3. SSL Configuration

The communication channel between application servers and `httpd` proxies uses HTTP by default. This channel can be secured using HTTPS by setting the `ssl` to `true`.

N.B. This HTTP/HTTPS channel should not be confused with the processing of HTTP/HTTPS client requests.

Attribute	Default	Description
<code>ssl</code>	<code>false</code>	Should connection to proxy use a secure socket layer
<code>sslCiphers</code>	The default JSSE cipher suites	Overrides the cipher suites used to init an SSL socket ignoring any unsupported ciphers
<code>sslProtocol</code>	TLS	Overrides the default SSL socket protocol.
<code>sslCertificateEncodingAlgorithm</code>	The default JSSE key manager algorithm	The algorithm of the key manager factory
<code>sslKeyStore</code>	<code>System.getProperty("user.home") + "/.keystore"</code>	The location of the key store containing client certificates
<code>sslKeyStorePass</code>	<code>changeit</code>	The password granting access to the key store
<code>sslKeyStoreType</code>	JKS	The type of key store
<code>sslKeyStoreProvider</code>	The default JSSE security provider	The key store provider
<code>sslTrustAlgorithm</code>	The default JSSE trust manager algorithm	The algorithm of the trust manager factory
<code>sslKeyAlias</code>		The alias of the key holding the client certificates in the key store
<code>sslCrIFile</code>		Certificate revocation list
<code>sslTrustMaxCertLength</code>	5	The maximum length of a certificate held in the trust store

Attribute	Default	Description
sslTrustStore	System.getProperty("javax.net.ssl.trustStore")	The trust store file containing the trust store
sslTrustStorePassword	System.getProperty("javax.net.ssl.trustStorePassword")	The password granting access to the trust store.
sslTrustStoreType	System.getProperty("javax.net.ssl.trustStoreType")	The trust store type
sslTrustStoreProvider	System.getProperty("javax.net.ssl.trustStoreProvider")	The trust store provider

8.1.4. HA Configuration

Additional configuration properties when mod_cluster is configured in clustered mode.

Attribute	Default	Description
masterPerDomain	false	If the domain directive is used, should HA partition use a singleton master per domain

8.1.5. Load Configuration

Additional configuration properties used when mod_cluster is configured in JBoss Web standalone or Tomcat.

Attribute	Default	Description
loadMetricClass	org.jboss.modcluster.load.metric.impl.BusyConnectorsLoadMetric	Class, BusyConnectorsLoadMetric object implementing org.jboss.load.metric.LoadMetric
loadMetricCapacity	1	The capacity of the load metric defined via the loadMetricClass property
loadHistory	9	The number of historic load values to consider in the load balance factor computation.
loadDecayFactor	2	The factor by which a historic load values should degrade in significance.

Load Metrics

9.1. Server-Side Load Metrics

A major feature of `mod_cluster` is the ability to use server-side load metrics to determine how best to balance requests.

The `DynamicLoadBalanceFactorProvider` bean computes the load balance factor of a node from a defined set of load metrics.

```
<bean name="DynamicLoadBalanceFactorProvider" class="org.jboss.modcluster.load.impl.DynamicLoadBalanceFactorProvider",exposedInterface=org.jboss.modcluster.load.impl.DynamicLoadBalanceFactorProviderMBean.class></annotation>

<constructor>

  <parameter>

    <set elementClass="org.jboss.modcluster.load.metric.LoadMetric">

      <inject bean="BusyConnectorsLoadMetric"/>

      <inject bean="HeapMemoryUsageLoadMetric"/>

    </set>

  </parameter>

</constructor>

<property name="history">9</property>

<property name="decayFactor">2</property>

</bean>
```

Load metrics can be configured with an associated weight and capacity:

1. The weight (default is 1) indicates the significance of a metric with respect to the other metrics. For example, a metric of weight 2 will have twice the impact on the overall load factor than a metric of weight 1.
2. The capacity of a metric serves 2 functions:

Each load metric contributes a value to the overall load factor of a node. The load factors from each metric are aggregated according to their weights.

In general, the load factor contribution of given metric is: (load / capacity) * weight / total weight.

The `DynamicLoadBalanceFactorProvider` applies a time decay function to the loads returned by each metric. The aggregate load, with respect to previous load values, can be expressed by the following formula:

$$L = (L_0 + L_1/D + L_2/D^2 + L_3/D^3 + \dots + L_H/D^H) * (1 + D + D^2 + D^3 + \dots + D^H)$$

... or more concisely as:

$$L = \sum_{i=0}^H L_i / D^i * \sum_{i=0}^H D^i$$

... where D = `decayFactor` and H = `history`.

Setting `history = 0` effectively disables the time decay function and only the current load for each metric will be considered in the load balance factor computation.

The `mod_cluster` proxy module expects the load factor to be an integer between 0 and 100, where 0 indicates max load and 100 indicates zero load. Therefore, the final load balance factor sent to the proxy = $100 - (L * 100)$.

While you are free to write your own load metrics, the following `LoadMetrics` are available out of the box:

9.2. Web Container metrics

1. `ActiveSessionsLoadMetric`

- * Requires an explicit capacity
- * Uses `SessionLoadMetricSource` to query session managers
- * Analogous to `method=S` in `mod_jk`

e.g.

```
<bean name="ActiveSessionsLoadMetric" class="org.jboss.modcluster.load.metric.impl.ActiveSessionsLoadMetric"
service="ActiveSessionsLoadMetric",exposedInterface=org.jboss.modcluster.load.metric.LoadMetricMBean.class></
annotation>

<constructor>
```

```

    <parameter><inject bean="SessionLoadMetricSource"/></parameter>

    </constructor>

    <property name="capacity">1000</property>

    </bean>

    <bean name="SessionLoadMetricSource" class="org.jboss.modcluster.load.metric.impl.SessionLoadMetricSource">

        <constructor>

            <parameter class="javax.management.MBeanServer"><inject bean="JMXKernel" property="mbeanServer"/></parameter>

            </constructor>

        </bean>

```

2. BusyConnectorsLoadMetric

* Returns the percentage of connector threads from the thread pool that are busy servicing requests

* Uses ThreadPoolLoadMetricSource to query connector thread pools

* Analogous to method=B in mod_jk

e.g.

```

    <bean name="BusyConnectorsLoadMetric" class="org.jboss.modcluster.load.metric.impl.BusyConnectorsLoadMetric"
        service="BusyConnectorsLoadMetric",exposedInterface=org.jboss.modcluster.load.metric.LoadMetricMBean.class)</
        annotation>

        <constructor>

            <parameter><inject bean="ThreadPoolLoadMetricSource"/></parameter>

        </constructor>

    </bean>

    <bean name="ThreadPoolLoadMetricSource" class="org.jboss.modcluster.load.metric.impl.ThreadPoolLoadMetricSource">

        <constructor>

            <parameter class="javax.management.MBeanServer"><inject bean="JMXKernel" property="mbeanServer"/></parameter>

```

```
</constructor>

</bean>
```

3. ReceiveTrafficLoadMetric

- * Returns the incoming request traffic in KB/sec
- * Requires an explicit capacity
- * Uses RequestProcessorLoadMetricSource to query request processors
- * Analogous to method=T in mod_jk

e.g.

```
<bean name="ReceiveTrafficLoadMetric" class="org.jboss.modcluster.load.metric.impl.ReceiveTrafficLoadMetric"
    service="ReceiveTrafficLoadMetric",exposedInterface=org.jboss.modcluster.load.metric.LoadMetricMBean.class></
annotation>

<constructor>

    <parameter class="org.jboss.modcluster.load.metric.impl.RequestProcessorLoadMetricSource" bean="RequestProcessorLoadMetricSource"/
    ></parameter>

</constructor>

<property name="capacity">1024</property>

</bean>

<bean name="RequestProcessorLoadMetricSource" class="org.jboss.modcluster.load.metric.impl.RequestProcessorLoadMetricSource">

    <constructor>

        <parameter class="javax.management.MBeanServer"><inject bean="JMXKernel" property="mbeanServer"/
        ></parameter>

    </constructor>

</bean>
```

4. SendTrafficLoadMetric

- * Returns the outgoing request traffic in KB/sec
- * Requires an explicit capacity

- * Uses RequestProcessorLoadMetricSource to query request processors

- * Analogous to method=T in mod_jk

e.g.

```
<bean name="SendTrafficLoadMetric" class="org.jboss.modcluster.load.metric.impl.SendTrafficLoadMetric" mod
web:service=SendTrafficLoadMetric",exposedInterface=org.jboss.modcluster.load.metric.LoadMetricMBean.class)</
annotation>

<constructor>

    <parameter>org.jboss.modcluster.load.metric.impl.RequestProcessorLoadMetricSource</parameter>

</constructor>

<property name="capacity">512</property>

</bean>
```

5. RequestCountLoadMetric

- * Returns the number of requests/sec

- * Requires an explicit capacity

- * Uses RequestProcessorLoadMetricSource to query request processors

- * Analogous to method=R in mod_jk

e.g.

```
<bean name="RequestCountLoadMetric" class="org.jboss.modcluster.load.metric.impl.RequestCountLoadMetric"
o:service=RequestCountLoadMetric",exposedInterface=org.jboss.modcluster.load.metric.LoadMetricMBean.class)</
annotation>

<constructor>

    <parameter>org.jboss.modcluster.load.metric.impl.RequestProcessorLoadMetricSource</parameter>

</constructor>

<property name="capacity">1000</property>
```

```
</bean>
```

9.3. System/JVM metrics

1. AverageSystemLoadMetric

- * Returns CPU load
- * Requires Java 1.6+.
- * Uses OperatingSystemLoadMetricSource to generically read attributes

e.g.

```
<bean name="AverageSystemLoadMetric" class="org.jboss.modcluster.load.metric.impl.AverageSystemLoadMetric" service="AverageSystemLoadMetric",exposedInterface=org.jboss.modcluster.load.metric.LoadMetricMBean.class)</bean>

<annotation>

<constructor>

<parameter><inject bean="OperatingSystemLoadMetricSource"/></parameter>

</constructor>

</bean>

<bean name="OperatingSystemLoadMetricSource" class="org.jboss.modcluster.load.metric.impl.OperatingSystemLoadMetricSource" service="OperatingSystemLoadMetricSource">

</bean>
```

2. SystemMemoryUsageLoadMetric

- * Returns system memory usage
- * Requires com.sun.management.OperatingSystemMXBean (available in Sun's JDK or OpenJDK)
- * Uses OperatingSystemLoadMetricSource to generically read attributes

e.g.

```
<bean name="SystemMemoryUsageLoadMetric" class="org.jboss.modcluster.load.metric.impl.SystemMemoryUsageLoadMetric" service="SystemMemoryUsageLoadMetric">
```

```

e=SystemMemoryUsageLoadMetric",exposedInterface=org.jboss.modcluster.load.metric.LoadMetricMBean.class)</
annotation>

<constructor>

<parameter><inject bean="OperatingSystemLoadMetricSource"/></parameter>

</constructor>

</bean>

```

3. HeapMemoryUsageLoadMetric

* Returns the heap memory usage as a percentage of max heap size

e.g.

```

ice=HeapMemoryUsageLoadMetric",exposedInterface=org.jboss.modcluster.load.metric.LoadMetricMBean.class)</
annotation>

</bean>

```

9.4. Other metrics

1. ConnectionPoolUsageLoadMetric

* Returns the percentage of connections from a connection pool that are in use.

* Uses ConnectionPoolLoadMetricSource to query JCA connection pools

e.g.

```

e=ConnectionPoolUsageLoadMetric",exposedInterface=org.jboss.modcluster.load.metric.LoadMetricMBean.class)</
annotation>

<constructor>

<parameter><inject bean="ConnectionPoolLoadMetricSource"/></parameter>

</constructor>

```

```
</bean>

<bean name="ConnectionPoolLoadMetricSource" class="org.jboss.modcluster.load.metric.impl.ConnectionPoolLoadMetricSource">

  <constructor>

    <parameter class="javax.management.MBeanServer"><inject bean="JMXKernel" property="mbeanServer"/></parameter>

  </constructor>

</bean>
```

Installation of the java part

10.1. Install the Java-side binaries

First, extract the java-side binary to a temporary directory. The following assumes it was extracted to /tmp/mod-cluster

Your next step depends on whether your target server is JBoss AS or JBossWeb/Tomcat.

10.1.1. Installing in JBoss AS 6.0.0.M1 and up

You don't need to do anything to install the java-side binaries in AS 6.x; it's part of the AS distribution's default, standard and all configurations.

10.1.2. Installing in JBoss AS 5.x

Assuming \$JBOSS_HOME indicates the root of your JBoss AS install and that you want to use mod_cluster in the AS's all config:

```
cp -r /tmp/mod-cluster/mod-cluster.sar $JBOSS_HOME/server/all/deploy
```

10.1.3. Installing in JBoss Web or Tomcat

Assuming \$CATALINA_HOME indicates the root of your JBossWeb or Tomcat install:

```
cp -r /tmp/mod-cluster/JBossWeb-Tomcat/lib/*  
$CATALINA_HOME/lib/
```

10.1.4. Installing in JBoss AS 4.2.x or 4.3.x

Assuming \$JBOSS_HOME indicates the root of your JBoss AS install and that you want to use mod_cluster in the AS's all config:

```
cp -r /tmp/mod-cluster/JBossWeb-Tomcat/lib/mod-cluster.jar $JBOSS_HOME/server/all/  
deploy/jboss-web.deployer
```

TODO

Using https between httpd and JBossWEB

The feature is only needed if the connection between httpd and JBossWEB is not secure it needs some more resources because crypto is used.

```
SSLProxyEngine On
SSLProxyVerify require
SSLProxyCACertificateFile conf/cacert.pem
SSLProxyMachineCertificateFile conf/proxy.pem
```

conf/proxy.pem should contain both key and certificate (and the certificate must be trusted by Tomcat).conf/cacert.pem must contain the CA that signed the Tomcat certificate.The <connector/> should use https.

Using SSL in mod_cluster

There are 2 connections between the cluster and the front-end. Both could be encrypted. That chapter describes how to encrypt both connections.

12.1. Using SSL between JBossWEB and httpd

As the ClusterListener allows to configure httpd it is advised to use SSL for that connection. The most easy is to use a virtual host that will only be used to receive information from JBossWEB. Both side need configuration.

12.1.1. Apache httpd configuration part

[mod_ssl](http://httpd.apache.org/docs/2.2/mod/mod_ssl.html) [http://httpd.apache.org/docs/2.2/mod/mod_ssl.html] of httpd is using to do that. See in one example how easy the configuration is:

```
Listen 6666
<VirtualHost 10.33.144.3:6666>
    SSLEngine on
    SSLCipherSuite AES128-SHA:ALL:!ADH:!LOW:!MD5:!SSLV2:!NULL
    SSLCertificateFile conf/server.crt
    SSLCertificateKeyFile conf/server.key
    SSLCACertificateFile conf/server-ca.crt
    SSLVerifyClient require
    SSLVerifyDepth 10
</VirtualHost>
```

The conf/server.crt file is the PEM-encoded Certificate file for the VirtualHost it must be signed by a Certificate Authority (CA) whose certificate is stored in the sslTrustStore of the ClusterListener parameter.

The conf/server.key file is the file containing the private key.

The conf/server-ca.crt file is the file containing the certificate of the CA that have signed the client certificate JBossWEB is using. That is the CA that have signed the certificate corresponding to the sslKeyAlias stored in the sslKeyStore of the ClusterListener parameters.

12.1.2. ClusterListener configuration part

There is a [wiki](http://www.jboss.org/community/docs/DOC-9300) [http://www.jboss.org/community/docs/DOC-9300] describing the SSL parameters of the ClusterListener. See in one example how easy the configuration is:

```
<Listener className="org.jboss.web.cluster.ClusterListener"
```

```
ssl="true"
sslKeyStorePass="changeit"
sslKeyStore="/home/jfclere/CERTS/CA/test.p12"
sslKeyStoreType="PKCS12"
sslTrustStore="/home/jfclere/CERTS/CA/ca.p12"
sslTrustStoreType="PKCS12" sslTrustStorePassword="changeit"
/>
```

The `sslKeyStore` file contains the private key and the signed certificate of the client certificate JBossWEB uses to connect to `httpd`. The certificate must be signed by a Certificate Authority (CA) who certificate is in the `conf/server-ca.crt` file of the `httpd`

The `sslTrustStore` file contains the CA certificate of the CA that signed the certificate contained in `conf/server.crt` file.

12.1.3. mod-cluster-jboss-beans configuration part

The `mod-cluster-jboss-beans.xml` in `$JBOSS_HOME/server/profile/deploy/mod-cluster.sar/META-INF` in the ClusterConfig you are using you should have something like:

```
<property name="ssl">true</property>
<property name="sslKeyStorePass">changeit</property>
<property name="sslKeyStore">/home/jfclere/CERTS/test.p12</property>
<property name="sslKeyStoreType">pkcs12</property>
<property name="sslTrustStore">/home/jfclere/CERTS/ca.p12</property>
<property name="sslTrustStoreType">pkcs12</property>
<property name="sslTrustStorePassword">changeit</property>
```

12.1.4. How the diferent files were created

The files were created using OpenSSL utilities see [OpenSSL](http://www.openssl.org/) [<http://www.openssl.org/>] `CA.pl` (`/etc/pki/tls/misc/CA` for example) has been used to create the test Certificate authority, the certificate requests and private keys as well as signing the certificate requests.

12.1.4.1. Create the CA

1. Create a work directory and work for there:

```
mkdir -p CERTS/Server
cd CERTS/Server
```

2. Create a new CA:

```
/etc/pki/tls/misc/CA -newca
```

That creates a directory for example ../../CA that contains a cacert.pem file which content have to be added to the conf/server-ca.crt described above.

12.1.4.2. Create the server certificate

1. Create a new request:

```
/etc/pki/tls/misc/CA -newreq
```

That creates 2 files named newreq.pem and newkey.pem. newkey.pem is the file conf/server.key described above.

2. Sign the request:

```
/etc/pki/tls/misc/CA -signreq
```

That creates a file named newcert.pem. newcert.pem is the file conf/server.crt described above. At that point you have created the SSL stuff needed for the VirtualHost in httpd. You should use a browser to test it after importing in the browser the content of the cacert.pem file.

12.1.4.3. Create the client certificate

1. Create a work directory and work for there:

```
mkdir -p CERTS/Client  
cd CERTS/Client
```

2. Create request and key for the JBossWEB part.

```
/etc/pki/tls/misc/CA -newreq
```

That creates 2 files: Request is in newreq.pem, private key is in newkey.pem

3. Sign the request.

```
/etc/pki/tls/misc/CA -signreq
```

That creates a file: newcert.pem

4. Don't use a passphrase when creating the client certificate or remove it before exporting:

```
openssl rsa -in newkey.pem -out key.txt.pem  
mv key.txt.pem newkey.pem
```

5. Export the client certificate and key into a p12 file.

```
openssl pkcs12 -export -inkey newkey.pem -in newcert.pem -out test.p12
```

That is the sslKeyStore file described above (/home/jfclere/CERTS/CA/test.p12)

12.2. Using SSL between httpd and JBossWEB

Using https allows to encrypt communications between httpd and JBossWEB. But due to the resources it needs that no advised to use it in high load configuration.

(See [Encrypting connection between httpd and TC](http://www.jboss.org/community/docs/DOC-9701) [http://www.jboss.org/community/docs/DOC-9701] for detailed instructions).

12.3. Forwarding SSL browser informations when using http/https between httpd and JBossWEB

When using http or https between httpd and JBossWEB you need to use the SSLValve and export the SSL variable as header in the request in httpd. If you are using AJP, mod_proxy_ajp will read the SSL variables and forward them to JBossWEB automatically.

(See [Forwarding SSL environment when using http/https proxy](http://www.jboss.org/community/docs/DOC-11988) [http://www.jboss.org/community/docs/DOC-11988] for detailed instructions).

The SSL variable used by mod_proxy_ajp are the following:

1. "HTTPS" SSL indicateur.
2. "SSL_CLIENT_CERT" Chain of client certificates.
3. "SSL_CIPHER" Cipher used.
4. "SSL_SESSION_ID" Id of the session.

Forwarding SSL browser informations when
using http/https between httpd and JBossWEB

5. "SSL_CIPHER_USEKEYSIZE" Size of the key used.

Migration from mod_jk

Mod_cluster only support Apache httpd, there are no plan to support IIS or Iplanet.

The migration from mod_jk to mod_cluster is not very complex. Only very few worker properties can't be mapped to mod_cluster parameters.

Here is the table of worker properties and how to transfer them in the ClusterListener parameters.

mod_jk worker property	ClusterListener parameter	Remarks
host	-	It is read from the <Connector/> Address information
port	-	It is read from the <Connector/> Port information
type	-	It is read from the <Connector/> Protocol information
route	-	It is read from the <Engine/> JVMRoute information
domain	domain	That is not supported in this version
redirect	-	The nodes with loadfactor = 0 are standby nodes they will be used no other nodes are available
socket_timeout	nodeTimeout	Default 10 seconds
socket_keepalive	-	KEEP_ALIVE os is always on in mod_cluster
connection_pool_size	-	The max size is calculated to be AP_MPMQ_MAX_THREADS+1 (max)
connection_pool_minsize	smax	The default is max
connection_pool_timeout	ttd	Time to live when over smax connections. The default is 60 seconds
-	workerTimeout	Max time to wait for a free worker default 1 second
retries	maxAttempts	Max retries before returning an error Default: 3
recovery_options	-	mod_cluster behave like mod_jk with value 7

mod_jk worker property	ClusterListener parameter	Remarks
fail_on_status	-	Not supported
max_packet_size	iobuffersize/receivebuffersize	Not supported in this version. Use ProxyIOBufferSize
max_reply_timeouts	-	Not supported
recover_time	-	The ClusterListener will tell (via a STATUS message) mod_cluster that the node is up again
activation	-	mod_cluster receives this information via ENABLE/DISABLE/STOP messages
distance	-	mod_cluster handles this via the loadfactor logic
mount	-	The context "mounted" automatically via the ENABLE-APP messages. ProxyPass could be used too
secret	-	Not supported
connect_timeout	-	Not supported. Use ProxyTimeout or server TimeOut (Default 300 seconds)
prepost_timeout	ping	Default 10 seconds
reply_timeout	-	Not supported. Use ProxyTimeout or server TimeOut? directive (Default 300 seconds)

Migration from mod_proxy

As mod_cluster is a sophisticated balancer migration from mod_proxy to mod_cluster is strait forward. mod_cluster replaces a reverse proxy with loadbalancing. A reversed proxy is configured like:

```
ProxyRequests Off
<Proxy *>
Order deny,allow
Allow from all
</Proxy>
ProxyPass /foo http://foo.example.com/bar
ProxyPassReverse /foo http://foo.example.com/bar
```

All the general proxy parameters could be used in mod_cluster they work like in mod_proxy, only the balancers and the workers definitions are slightly different.

14.1. Workers

Mod_proxy Parameter	ClusterListener parameter	Remarks
min	-	Not supported in this version
max	-	mod_cluster uses mod_proxy default value
smax	smax	Same as mod_proxy
ttd	ttd	Same as mod_proxy
acquire	workerTimeout	Same as mod_proxy acquire but in seconds
disablereuse	-	mod_cluster will disable the node in case of error and the ClusterListener will for the reuse via the STATUS message
flushPackets	flushPackets	Same as mod_proxy
flushwait	flushwait	Same as mod_proxy
keepalive	-	Always on: OS KEEP_ALIVE is always used. Use connectionTimeout in the <Connector> if needed
lbset	-	Not supported

Mod_proxy Parameter	ClusterListener parameter	Remarks
ping	ping	Same as mod_proxy Default value 10 seconds
lbfactor	-	The load factor is received by mod_cluster from a calculated value in the ClusterListener
redirect	-	Not supported lbfactor sent to 0 makes a standby node
retry	-	ClusterListener will test when the node is back online
route	JVMRoute	In fact JBossWEB via the JVMRoute in the Engine will add it
status	-	mod_cluster has a finer status handling: by context via the ENABLE/STOP/DISABLE/REMOVE application messages. hot-standby is done by lbfactor = 0 and Error by lbfactor = 1 both values are sent in STATUS message by the ClusterListener
timeout	nodeTimeout	Default wait for ever (http://httpd.apache.org/docs/2.2/mod/mod_proxy.html is wrong there)
ttl	ttl	Default 60 seconds

14.2. Balancers

Mod_proxy Parameter	ClusterListener parameter	Remarks
lbmethod	-	There is only one load balancing method in mod_cluster "cluster_byrequests"
maxattempts	maxAttempts	Default 3
nofailover	stickySessionForce	Same as in mod_proxy
stickysession	StickySessionCookie/ StickySessionPath	The 2 parameters in the ClusterListener are combined in one that behaves like in mod_proxy

Mod_proxy Parameter	ClusterListener parameter	Remarks
timeout	workerTimeout	Default 1 seconds

Load Balancing Demo Application

15.1. Overview

The mod_cluster distribution includes a demo application that helps demonstrate how different server-side scenarios affect the routing of client requests by the load balancer. The demo application is located in the mod_cluster distribution's demo directory.

The demo application consists of two components:

1. The first component is a war file that needs to be deployed in JBossWeb/Tomcat/JBoss AS. The war includes a number of servlets:
2. The second component is a GUI application that allows a user to launch a pool of threads that repeatedly make requests to the load balancer. The requests are ultimately routed to the demo war's primary servlet. The application tracks which servers are handling the requests and displays this information in a chart.

The application can also send separate requests to the demo war's load generation servlets, allowing the user to see how different load conditions affect the balancing of requests.

Note that the demo application does not actually depend on mod_cluster in any way. Its only dependency is on JBossWeb/Tomcat.¹ So, the demo could be used to demonstrate the effect of different server-side scenarios on the routing decisions made by any load balancer, including mod_jk, mod_proxy or the various hardware load balancers.

Note also that this demo application is not intended to be used as a load testing tool; i.e. something that can demonstrate the maximum load a cluster configuration can handle. Using it as such has a good chance of showing you the maximum load the client can generate rather than the maximum load your cluster can handle.

15.2. Basic Usage

To run the demo application:

1. Unpack the mod_cluster distribution on your filesystem. Here we assume it has been unzipped to /home/bes/mod_cluster or C:\mod_cluster.
2. Install mod_cluster into your httpd server as described at [Installation of the httpd part](#)
3. Install mod_cluster into your JBossAS/JBossWeb/Tomcat servers as described at [Installation on the Java side](#)

¹The demo's "Datasource Use" load generation scenario requires the use of JBoss Application Server.

4. Start httpd and your JBossAS/JBossWeb/Tomcat servers
5. Deploy the load-demo.war found in the distribution's demo/server folder to your JBossAS/JBossWeb/Tomcat servers.
6. Start the demo application:
 - a. On *nix:

```
[bes@besdev ~]$ cd /home/bes/mod_cluster/demo/client  
[bes@besdev mod_cluster]$ ./run-demo.sh
```

- b. On Windows:

```
C:\>cd mod_cluster\demo\client  
C:\mod_cluster\demo\client>run-demo
```

The screenshot shows a web application titled "Load Balancing Demonstration". It has four tabs: "Content Control", "Server Load Control", "Request Balancing", and "Session Balancing". The "Server Load Control" tab is currently selected. Below the tabs, there are four input fields and a button:

- Target Hostname:** A text input field containing "localhost".
- Target Port:** A text input field containing "8000".
- Load Creation Action:** A dropdown menu with "Connection Pool Use" selected.
- Number of Connections:** A text input field containing "50".
- Duration:** A text input field containing "15".
- Create Load:** A blue button with white text.

7. Configure the hostname and address of the httpd server, the number of client threads, etc and click the "Start" button. See [Client Driver Configuration Options](#) for details on the configuration options.
8. Switch to the "Request Balancing" tab to see how many requests are going to each of your JBossAS/JBossWeb/Tomcat servers:
9. Switch to the "Session Balancing" tab to see how many active sessions² are being hosted by each of your JBossAS/JBossWeb/Tomcat servers:
10. Stop some of your JBossAS/JBossWeb/Tomcat servers and/or undeploy the load-demo.war from some of the servers and see the effect this has on load balancing.
11. Restart some of your JBossAS/JBossWeb/Tomcat servers and/or re-deploy the load-demo.war to some of the servers and see the effect this has on load balancing.

12Experiment with adding artificial load to one or more servers to see what effect that has on load balancing. See [Load Generation Scenarios](#) for details.

Most of the various panels in application interface also present information on the current status on any client threads. "Total Clients" is the number of client threads created since the last time the "Start" button was pushed. "Live Clients" is the number of threads currently running. "Failed Clients" is the number of clients that terminated abnormally; i.e. made a request that resulted in something other than an HTTP 200 response.

15.3. Client Driver Configuration Options

The configuration of the client is driver is done via the application's "Client Control" tab:

The screenshot shows the 'Load Balancing Demonstration' application window. The 'Client Control' tab is selected, showing the following configuration options:

Configuration Option	Value
Target Hostname:	localhost
Target Port:	8000
Load Creation Action:	Connection Pool Use
Number of Connections:	50
Duration:	15

A 'Create Load' button is located below the configuration fields.

The panel includes the following options:

1. Proxy Hostname: Hostname of the load balancer or the IP address on which it is listening for requests³
2. Proxy Port: Port on which the load balancer is listening for requests⁴
3. Context Path: Portion of the request URL that specifies the request is for the load-demo.war
4. Session Life: Number of seconds a client thread should use a session before invalidating or abandoning it. Generally it is good to keep this to a small value; otherwise the use of session stickiness will prevent changes in server load from affecting the load balancer's routing decisions. With sticky sessions enabled (strongly recommended), it is the creation of a new session that allows the load balancer to try to balance load.
5. Invalidate: Controls what the client thread should do when it stops using a session because Session Life has passed. If checked, the driver will send a request that results in the session being invalidated. If unchecked, the session will just be abandoned, and will continue to exist on the server until Session Timeout seconds have passed. (In the future this will likely be changed to a percentage input, so X% can be invalidated, the rest abandoned.)
6. Session Timeout: Number of seconds a session can remain unused before the server is free to expire it. Unchecking Invalidate and setting a high value relative to Session Life allows a significant number of unused sessions to accumulate on the server.
7. Num Threads: Number of client threads to launch. Each thread repeatedly makes requests until the "Stop" button is pushed or a request receives a response other than HTTP 200.
8. Sleep Time: Number of ms the client threads should sleep between requests.
9. Startup Time: Number of seconds over which the application should stagger the start of the client threads. Staggering the start is advised as it avoids the unnatural situation where for the life of the demonstration all sessions start at about the same time and then are invalidated or abandoned at the same time. Staggering the start allows the load balancer to continually see new sessions and decide how to route them.

15.4. Load Generation Scenarios

You can use the application's GUI to instruct individual servers to artificially generate various types of load, and then track how that load affects request and session balancing. Load generation is controlled via the application's "Server Load Control" tab:

The screenshot shows a web application titled "Load Balancing Demonstration". It has four tabs: "Request Control", "Server Load Control" (which is selected), "Request Balancing", and "Session Balancing". Below the tabs, there are several input fields and a button. The "Target Hostname" field contains "localhost". The "Target Port" field contains "8000". The "Load Creation Action" is a dropdown menu currently showing "Connection Pool Use". The "Number of Connections" field contains "50". The "Duration" field contains "15". At the bottom center is a button labeled "Create Load".

Load Balancing Demonstration			
Request Control	Server Load Control	Request Balancing	Session Balancing
Target Hostname:	<input type="text" value="localhost"/>	Target Port:	<input type="text" value="8000"/>
Load Creation Action:	<input style="border: 1px solid #ccc;" type="text" value="Connection Pool Use"/>		
Number of Connections	<input type="text" value="50"/>		
Duration	<input type="text" value="15"/>		
<input type="button" value="Create Load"/>			

The panel includes the following options:

1. Target Hostname and Target Port: The hostname or IP address of the server on which you want load generated. There are two strategies for setting these:
 - a. You can use the hostname and port of the load balancer, in which case the load balancer will pick a backend server and route the request to it. Note the client application does not maintain a session cookie for these requests, so if you invoke another server load generation request, you shouldn't expect the same server to handle it.
 - b. If the JBoss AS/JBossWeb/Tomcat servers are running the `HttpConnector` as well as the `AJP` connector, you can specify the address and port on which a particular server's `HttpConnector` is listening. The standard port is 8080.
2. Load Creation Action: Specifies the type of load the target server should generate. See below for details on the available load types.








3. Params: Zero or more parameters to pass to the specified load creation servlet. For example, in the screenshot above, Number of Connections and Duration. How many parameters are displayed, their name and their meaning depend on the selected Load Creation Action. The label for each parameter includes a tooltip that explains its use.

The available Load Creation Actions are as follows:





1. Active Sessions: Generates server load by causing session creation on the target server.
2. Datasource Use: Generates server load by taking connections from the java:DefaultDS datasource for a period¹
3. Connection Pool Use: Generates server load by tying up threads in the webserver connections pool for a period
4. Heap Memory Pool Use: Generates server load by filling 50% of free heap memory for a period
5. CPU Use: Generates server CPU load by initiating a tight loop in a thread
6. Server Receive Traffic: Generates server traffic receipt load by POSTing a large byte array to the server once per second for a period
7. Server Send Traffic: Generates server traffic send load by making a request once per second to which the server responds with a large byte array
8. Request Count: Generates server load by making numerous requests, increasing the request count on the target server.

Change Log


16.1. 1.0.10 (14 April 2011)

	Apache with mod_cluster refuses to start at first, after 7 retries it starts up (MODCLUSTER-120) (jfclere)
	NPE when overriding default load metric in ModClusterListener (MODCLUSTER-183) (jfclere)
	Only 100 webapps are detected : Maxcontext doesn't seem to be read (MODCLUSTER-221) (jfclere)
	when using tomcat manager webapp to stop/start an application the start is ignored by mod_cluster (MODCLUSTER-224) (jfclere)
	Tomcat6 throws java.lang.NoSuchMethodError instead a localised error messages. (MODCLUSTER-228) (jfclere)
	Add a note explaining where Maxcontext must be put in httpd.conf and error when it is in wrong place. (MODCLUSTER-225) (jfclere)
	Make mod_cluster manager tolerant to F5 page refresh when disabled context. (MODCLUSTER-218) (jfclere)




16.2. 1.0.9 (17 March 2011)

	Query strings are not forwarded to nodes. (MODCLUSTER-118) (jfclere)
	memory usage growing. (MODCLUSTER-214) (jfclere)
	mod_rewrite PT doesn't work. (MODCLUSTER-213) (jfclere)
	UseAlias broken. (MODCLUSTER-212) (jfclere)





16.3. 1.0.8 (15 Feb 2011)

	mod_cluster failover does not work for a / webappcontext when the / root context exists REOPENED. (MODCLUSTER-188) (jfclere)
---	--

16.4. 1.0.7 (not released)



	STATUS should retry the worker even if there was an error before. (MODCLUSTER-133) (jfclere)
	httpd cores after graceful restart after the mod_cluster configuration is added (MODCLUSTER-206) (jfclere)
	mod_cluster 1.0.0 docs step 12.1.4.3. is wrong (MODCLUSTER-193) (jfclere)

16.5. 1.0.6 (05 January 2011)












	Regressions caused by MODCLUSTER-190 (MODCLUSTER-190) (jfclere)
	Default maxAttempts set to -1 instead of 1 (MODCLUSTER-200) (jfclere)
	OK STATUS MCMP messages are send before the connector is started (MODCLUSTER-202) (jfclere)
	Quotes in jsessionId causing sticky sessions to fail (MODCLUSTER-203) (jfclere)

16.6. 1.0.5 (23 November 2010)


	Alias max length too low (MODCLUSTER-173) (jfclere)
	Advertise not configured error message in log is actually a warning (MODCLUSTER-176) (jfclere)
	mod_cluster failover does not work for a / webappcontext when the / root context exists (MODCLUSTER-188) (jfclere)
	mod_cluster issues an ENABLE-APP too early in the webapp lifecycle (MODCLUSTER-190) (jfclere)






	mod_cluster-manager does not update when a node is taken down (MODCLUSTER-195) (jfclere)
	Incorrect routing of requests when one context root is the prefix of another (MODCLUSTER-196) (jfclere)

16.7. 1.0.4 (27 July 2010)










	Separate 1.0.x doc from 1.1.x (jfclere)
	Advertise does not work on name-based virtual hosts (MODCLUSTER-122) (jfclere)
	Microcontainer does not choose the right constructor when creating RequestCountLoadMetric (MODCLUSTER-126) (pferraro)
	proxy_cluster_isup: Can't find worker for 2 (just after CONFIG). (MODCLUSTER-149) (jfclere)
	SystemMemoryUsageLoadMetric returns wrong load metric. (MODCLUSTER-157) (pferraro)
	ManagerBalancerName doesn't work. (MODCLUSTER-163) (jfclere)
	Fix class name of MBeanAttributeRatioLoadMetric in MC config (MODCLUSTER-174) (pferraro)
	Wrong configuration could cause a core (MODCLUSTER-175) (jfclere)
	ping and nodeTimeout interact. (MODCLUSTER-180) (jfclere)
	Add a solaris10 64 bits sparc in the bundles. (MODCLUSTER-137) (jfclere)
	Avoid unnecessary open sockets for non-master nodes (MODCLUSTER-158) (pferraro)


16.8. 1.0.3.GA (15 January 2010)

	HAModClusterService can throw NullPointerExceptions after a partition merge. (MODCLUSTER-85) (pferraro)
---	---










	random 404 errors on increasing load. (MODCLUSTER-102) (jfclere)
	random 404 errors on increasing load. (MODCLUSTER-102) (jfclere)
	Advertise security key verification does not work. (MODCLUSTER-104) (jfclere)
	load-demo.war specifies non-existent servlet in web.xml. (MODCLUSTER-113) (pferraro)
	AdvertiseBindAddress does not default to the 23364 port. (MODCLUSTER-119) (jfclere)
	Mod_cluster does support more that 3 Alias in Host. (MODCLUSTER-121) (jfclere)
	Advertise does not work on name-based virtual hosts. (MODCLUSTER-122) (jfclere)

16.9. 1.0.2.GA (20 August 2009)




	Update httpd to 2.2.13. (MODCLUSTER-93) (jfclere)
	DISABLED applications act as STOPPED. (MODCLUSTER-96) (jfclere)
	getProxyInfo failed when there are too many nodes. (MODCLUSTER-94) (jfclere)
	mod_cluster_manager displays zeros instead values on Linux 32 bits. (MODCLUSTER-98) (jfclere)
	mod_cluster_manager doesn't seem to ENABLE/DISABLE the right context. (MODCLUSTER-99) (jfclere)
	Load balancing logic doesn't allow manual demo of load-balancing: it seems to go always to the first nodes. (MODCLUSTER-100) (jfclere)
	Alias from webapps/jboss-web.xml are not handled correctly in mod_cluster. (MODCLUSTER-89) (jfclere)
	mod_cluster-manager display corrupted when jboss is starting. (MODCLUSTER-95) (jfclere)
	


	ClassCastException upon redeploy after mod-cluster-jboss-beans.xml modification. (MODCLUSTER-88) (pferraro)
	admin-console should be in the excludedContexts. (MODCLUSTER-87) (pferraro)

16.10. 1.0.1.GA (29 July 2009)







	Fail-over from cleanly shutdown nodes ignores the domain information (MODCLUSTER-82) (jfclere)
	Add logic to support Alias (in server.xml) and NameVirtualHost (in httpd.conf) (MODCLUSTER-77) (jfclere)
	mod_cluster-manager: disabling a worker / domain (MODCLUSTER-50) (jfclere)
	mod_cluster-manager: disabling a worker / domain (MODCLUSTER-50) (jfclere)
	Allow static content to be served from proxy's local filesystem. (MODCLUSTER-37) (jfclere)
	httpd crashes when you try to reach an application disabled via /mod_cluster-manager. (MODCLUSTER-80) (jfclere)
	Disabled application gets request without sessionid. (MODCLUSTER-81) (jfclere)
	Add support for JBossAS 4.2.x (documentation) (MODCLUSTER-78) (jfclere)
	Add INFO and DUMP in the mod_cluster-manager. Partialy MODCLUSTER-79 (jfclere)

16.11. 1.0.0.GA (29 May 2009)







	Allow to use mod_balancer syntax like ProxyPass / balancer://mycluster in httpd VirtualHosts (MODCLUSTER-75) (jfclere)
	Add AdvertiseBindAddress to specify the bind address of the Advertise logic on httpd (MODCLUSTER-76) (jfclere)
	




	Fixed serialization regression with some rpc response content when in clustered mode. (pferraro)
	Various fixes to load simulation section of demo application. (pferraro)

16.12. 1.0.0.CR2 (29 April 2009)













	Add ability to disable/enable all/individual apps from mod_cluster-manager status page (MODCLUSTER-69) (jfclere)
	Add ability to disable/enable individual webapp via JMX (MODCLUSTER-68) (pferraro)
	openssl doesn't build on some platforms, update to 0.9.8k (MODCLUSTER-70) (jfclere)
	httpd crashes on windows when you kill JBoss AS (MODCLUSTER-67)
	Advertise worker thread can throw NullPointerException (MODCLUSTER-73) (pferraro)
	Simplify rpc response objects (pferraro)

16.13. 1.0.0.CR1 (24 March 2009)

	Add Bundle with only the modules libraries (MODCLUSTER-28) (jfclere)
	Windows bundles broken (MODCLUSTER-58) (jfclere)
	Mod_cluster-manager status page reports wrong information (MODCLUSTER-64) (jfclere)
	Undeploy an app or shut down server, clients with an existing session do not fail over. (MODCLUSTER-63) (jfclere)
	httpd.worker is a shell script (MODCLUSTER-56) (jfclere)
	update compoment bundle (httpd-2.2.11, openssl-0.9.8j and jk-1.2.27) (MODCLUSTER-55) (jfclere)










	HAModClusterService unnecessarily updates DRM each status cycle even if nothing changed (MODCLUSTER-60) (pferraro)
	Fixed heap usage load metric when max heap is undefined. (MODCLUSTER-62) (pferraro)
	Fixed advertise listener's pause/resume logic.

16.14. 1.0.0.Beta4 (14 Feb 2009)






	Remove incompatible mod_proxy_balancer from bundles (MODCLUSTER-46) (jfclere)
	Rename mod_sharedmem.so to mod_slotmem.so in win32/64 bundles (MODCLUSTER-45) (jfclere)
	Prevent Advertise to send bad X-Manager-Address when miss configured (MODCLUSTER-49) (jfclere)
	Prevent Advertise to send bad X-Manager-Address when miss configured (MODCLUSTER-49) (jfclere)
	Fix Advertise on win32/64 (MODCLUSTER-41) (mturk)
	Error state recording in ClusteredMCMPhandler can cause unnecessary reset requests (MODCLUSTER-53) (pferraro)
	mod_manager core on hp-ux (MODCLUSTER-43) (jfclere)
	Add missing mod_proxy_cluster.so in hp-ux (MODCLUSTER-42) (jfclere)
	Remove request session information when the response contains another session (MODCLUSTER-44) (jfclere)
	Fix AJP connector detection (TC6 jk connector weren't detected) (MODCLUSTER-52) (pferraro/jfclere)
	Forward unparsed uri to the node (instead a parsed one without query string) (jfclere)
	





	Master node enters tight loop calling getClusterCoordinatorState on remote node (MODCLUSTER-47) (pferraro)
--	--

16.15. 1.0.0.Beta3 (31 Jan 2009)


	Add a handler (mod_cluster-manager) to display node status in httpd (MODCLUSTER-27) (jfclere)
	exclude specific webapps from mod_cluster httpd registration (MODCLUSTER-22) (pferraro)
	change the default value for Maxnode to 20 (jfclere)
	Add a system property for specifying the proxyList (MODCLUSTER-36) (pferraro)
	fix the vhost corruption (MODCLUSTER-31) (jfclere)
	fix the win32/win64 build for the missing *.so files (MODCLUSTER-40) (jfclere)
	Allow empty proxyList (MODCLUSTER-38) (pferraro)
	advertisePort in {HA}ModClusterService default is wrong (MODCLUSTER-30) (jfclere)
	ActiveSessionsLoadMetric throws AttributeNotFoundException with distributable application (MODCLUSTER-29) (pferraro)

16.16. 1.0.0.Beta2 (12 Dec 2008)

	update httpd to 2.2.10. (jfclere)
	update openssl to 0.9.8i. (jfclere)
	Arrange mpm worker reslist logic. (jfclere)
	Improve integration of the test with the httpd. (jfclere)
	Fixed shutdown event handling. (pferraro)

	Separated configuration for clustered and non-clustered mode. (pferraro)
	Change the retry logic when sending to httpd: resend the message when read failed. (jfcclere)
	Refactored JBoss Web integration points (listeners/services). (pferraro)
	Added support for dynamic load balancing in JBoss Web standalone/Tomcat. (pferraro)

16.17. 1.0.0.Beta1 (05 Nov 2008)

	Initial beta release for use with JBossAS-5.0.x, JBossWEB-2.1.x and Tomcat-6.0.x.
---	---

Frequently Asked questions

17.1. What is Advertise

Advertise allows autodiscovery of httpd proxies by the cluster nodes. It is done by sending multicast messages from httpd to the cluster. The httpd specialized module: mod_advertise sends UDP message on a multicast group, both mod_advertise and the cluster listener joined the multicast group and the cluster receives the messages.

Example for mod_advertise message:

HTTP/1.0 200 OK

Date: Wed, 08 Apr 2009 12:26:32 GMT

Sequence: 16

Digest: f2d5f806a53effa6c67973d2ddcdd233

Server: 1b60092e-76f3-49fd-9f99-a51c69c89e2d

X-Manager-Address: 127.0.0.1:6666

X-Manager-Url: /bla

X-Manager-Protocol: http

X-Manager-Host: 10.33.144.3

The X-Manager-Address header value is used by the cluster logic to send information about the cluster to the proxy. It is the IP and port of the VirtualHost where mod_advertise is configured or URL parameter of the ServerAdvertise directive.

See [Proxy Discovery](#) in Configuration Properties and [mod_advertise](#) in Apache httpd configuration.

17.2. What to do if I don't want to use Advertise (multicast):

In the VirtualHost receiving the MCPM of httpd.conf don't use any Advertise directive or use:

ServerAdvertise Off

In mod_cluster-jboss-beans.xml add the addresses and ports of the VirtualHost to the proxyList property and set advertise to false, for example:

```
<property name="proxyList">10.33.144.3:6666,10.33.144.1:6666</property>
```

```
<property name="advertise">false</property>
```

17.3. I am using Tomcat 7 / 6 what should I do:

See at the end of [java configuration](#) You can't use the mod_cluster clustered mode with Tomcat so you get a loadbalancing logic similar to mod_jk but with a dynamic configuration.